

K-Nearest Neighbors (KNN)

A Step-by-Step, Practical Guide

This guide walks you through how KNN works from intuition to implementation, with a worked example, tips for choosing hyperparameters, and ready-to-use code.

What is KNN?

K-Nearest Neighbors is a simple, non-parametric, instance-based learning algorithm. To predict for a new sample, it looks at the k closest training points (its neighbors) using a distance metric, then **votes** (classification) or **averages** (regression) their targets. No explicit training step is required: the model is essentially the stored dataset.

When should you use KNN?

- You want a strong baseline with minimal tuning.
- Decision boundary may be irregular and non-linear.
- Small to medium datasets (thousands to low millions) and moderate feature counts.
- You can afford slower prediction time or use accelerated nearest-neighbor search.

KNN: Step-by-Step

Training

- Store the training data X and labels/targets y .
- Optionally scale features (recommended).

Prediction for a new sample x

- 1 Choose k (number of neighbors).
- 2 Choose a distance metric (e.g., Euclidean, Manhattan, Minkowski, cosine).
- 3 Compute distances from x to all training points.
- 4 Find the k points with the smallest distance.
- 5 **Classification**: majority vote (optionally weighted by $1/\text{distance}$).
- 6 **Regression**: average the targets (optionally distance-weighted).
- 7 Return the predicted class/target.

Common Distance Metrics

- **Euclidean** (L2): $d(x, z) = \sqrt{\sum (x_i - z_i)^2}$.
- **Manhattan** (L1): $d(x, z) = \sum |x_i - z_i|$.
- **Minkowski**: $d(x, z) = (\sum |x_i - z_i|^p)^{1/p}$, with $p \geq 1$.
- **Cosine**: $1 - (x \cdot z) / (||x|| ||z||)$, useful for high-dimensional sparse vectors.

Tip: Scale features (e.g., StandardScaler or MinMax) so that units don't distort distance.

Choosing k

- Use cross-validation to tune k. Odd k helps avoid ties in binary classification.
- Small k → low bias, high variance (risk of overfitting).
- Large k → higher bias, lower variance (can oversmooth).
- Heuristic: start with $k \approx \sqrt{n}$ and search around it.

Weighted KNN

Instead of equal votes, weight each neighbor by an inverse distance, e.g., $w = 1 / (d + \epsilon)$. This can improve performance when relevant neighbors are very close.

Complexity & Data Structures

- **Training:** $O(1)$ — just store the data (and maybe scale).
- **Prediction:** $O(n \cdot d)$ per query for naive search (n samples, d features).
- Speedups: KD-Tree / Ball-Tree for moderate d; approximate nearest neighbor (ANN) for high d / large n.

Practical Considerations

- **Categorical features:** encode (one-hot/ordinal); consider Hamming for binary.
- **Missing values:** impute before distance computation.
- **Class imbalance:** use class weights, stratified CV, or resampling.
- **Curse of dimensionality:** reduce d (PCA, feature selection) or use ANN methods.

Worked Example (Classification)

We have 2D points with classes A/B. Predict the class of $x = (3, 4)$ using Euclidean distance and $k = 3$.

Point	(x, y)	Class	Distance to $x_0=(3,4)$
P1	(2, 3)	A	1.414
P2	(5, 4)	B	2.000
P3	(3, 2)	A	2.000
P4	(7, 6)	B	4.472
P5	(1, 5)	A	2.236
P6	(4, 5)	B	1.414

Nearest 3 neighbors: P1 (A), P6 (B), P2 (B). Majority vote → **Class B**.

Pseudocode

```
KNN-Predict( $x_0$ , X, y, k, metric): # X: nxd matrix, y: length-n labels/targets # metric:
function to compute distance between two d-vectors
D = []
for i in 1..n: d = metric( $x_0$ , X[i])
D.append((d, y[i]))
sort D by distance ascending
N = first k elements of D
if classification: return majority_label(N) # or weighted by  $1/(d+\epsilon)$ 
else: # regression
return average_target(N) # or weighted average
```

Scikit-learn Example (Classification)

```
from sklearn.preprocessing import StandardScaler from sklearn.neighbors import
KNeighborsClassifier from sklearn.pipeline import make_pipeline from
sklearn.model_selection import cross_val_score pipe = make_pipeline(StandardScaler(),
KNeighborsClassifier(n_neighbors=5, weights='distance', metric='minkowski', p=2))
scores = cross_val_score(pipe, X, y, cv=5) print(scores.mean())
```

Pros & Cons

- **Pros:** simple, no training time, handles non-linear boundaries, competitive baseline.
- **Cons:** slow predictions on large n, sensitive to scale & noise, suffers in high dimensions.

Quick Checklist

- Scale features.
- Tune k with cross-validation (start near \sqrt{n}).
- Try weights='distance'.
- Pick metric suited to data; consider dimensionality reduction.
- For speed, use KD/Ball Trees (moderate d) or ANN methods (high d/large n).

You're set! Use this as a reference when building and tuning KNN models.